

Autoregressive moving average models

Kevin Kotzé

Contents

1. Introduction
2. Moving Average Models
3. Autoregressive Models
4. ARMA Models
5. Seasonal ARMA Models
6. Model Specification and Parameter Estimation
7. Structural Breaks
8. Conclusion

Univariate models for persistent data

- Dominant feature of many time series is that today's values are close to tomorrow's values
- Observations are not independent, but autocorrelated
- Need to account for this behaviour in the explained part of the model, otherwise it will be captured by the error, which violates the assumptions of the model
- Example of stochastic process:

$$y_t = 0.7y_{t-1} + \varepsilon_t$$

- This could represent an example of a linear stochastic difference equation, that includes discrete information
- Descriptive information should be used to populate the coefficient and random noise should be contained in the error

Moving average models

- Linear combination of white noise (i.e. ε_t), such that the $MA(1)$ may take the form,

$$y_t = \mu + \varepsilon_t + \theta\varepsilon_{t-1}$$

- where μ is a constant, while ε_t and ε_{t-1} are independent and identically distributed white noise, $\varepsilon_t \sim \text{i.i.d. } \mathcal{N}(0, \sigma^2)$
- To determine whether the $MA(1)$ process is stationary, we calculate the different moments

MA models - Expected Mean

- Note that $\mathbb{E}[\varepsilon_t] = 0$ and $\mathbb{E}[\varepsilon_t^2] = \sigma^2$,

$$\begin{aligned}\mathbb{E}[y_t] &= \mathbb{E}[\mu + \varepsilon_t + \theta\varepsilon_{t-1}] \\ &= \mu + \mathbb{E}[\varepsilon_t] + \theta\mathbb{E}[\varepsilon_{t-1}] \\ &= \mu\end{aligned}$$

- Since error terms are i. i. d. and their expected mean value is zero
- Hence, the mean for this process is μ , which is constant and does not depend on time

MA models - Variance

$$\begin{aligned}\text{var}[y_t] &= \mathbb{E}[y_t - \mathbb{E}[y_t]]^2 \\ &= \mathbb{E}[(\mu + \varepsilon_t + \theta\varepsilon_{t-1}) - \mu]^2 \\ &= \mathbb{E}[\varepsilon_t]^2 + 2\theta\mathbb{E}[\varepsilon_t\varepsilon_{t-1}] + \mathbb{E}[\theta\varepsilon_{t-1}]^2 \\ &= \sigma^2 + 0 + \theta\sigma^2 \\ &= (1 + \theta^2)\sigma^2\end{aligned}$$

- which is constant and does not depend on time

MA models - Covariance

- For the first lag,

$$\begin{aligned}\text{cov}[y_t, y_{t-1}] &= \mathbb{E} \left[(y_t - \mathbb{E}[y_t]) (y_{t-1} - \mathbb{E}[y_{t-1}]) \right] \\ &= \mathbb{E} \left[(\varepsilon_t + \theta \varepsilon_{t-1}) (\varepsilon_{t-1} + \theta \varepsilon_{t-2}) \right] \\ &= \mathbb{E} [\varepsilon_t \varepsilon_{t-1}] + \theta \mathbb{E} [\varepsilon_{t-1}^2] + \mathbb{E} [\theta \varepsilon_t \varepsilon_{t-2}] + \mathbb{E} [\theta^2 \varepsilon_{t-1} \varepsilon_{t-2}] \\ &= 0 + \theta \sigma^2 + 0 + 0 \\ &= \theta \sigma^2\end{aligned}$$

- which is constant and does not depend on time

MA models - Covariance

- For the general case of j lags,

$$\begin{aligned}\text{cov}[y_t, y_{t-j}] &= \mathbb{E} \left[(y_t - \mathbb{E}[y_t]) (y_{t-j} - \mathbb{E}[y_{t-j}]) \right] \\ &= \mathbb{E} \left[(\varepsilon_t + \theta \varepsilon_{t-1}) (\varepsilon_{t-j} + \theta \varepsilon_{t-j-1}) \right] \\ &= 0 \quad \text{for } j > 1\end{aligned}$$

- which is constant and does not depend on time

MA models - Stationarity

- Neither the mean, variance nor covariances depend on time
- Hence the $MA(1)$ process is covariance stationary
- Such a $MA(1)$ process is stationary regardless of the value θ

MA models - ACFs

- ACF for a $MA(1)$ may then be derived from the expression,

$$\rho(j) \equiv \frac{\gamma(j)}{\gamma(0)} = \frac{\text{cov}[y_t, y_{t-j}]}{\text{var}[y_t]}$$

- Hence,

$$\begin{aligned}\rho(1) &= \frac{\theta}{(1 + \theta^2)} \\ \rho(j) &= 0 \quad \text{for } j > 1\end{aligned}$$

- for lag orders $j > 1$, the autocorrelations are zero

Figure 1: Simulated $MA(1)$: $\varepsilon_t - 0.5\varepsilon_{t-1}$

Figure 2: Autocorrelation Functions for $MA(1)$: $\varepsilon_t - 0.5\varepsilon_{t-1}$

MA models - Higher Order

- Finite order $MA(q)$ process may be,

$$y_t = \mu + \varepsilon_t + \theta_1 \varepsilon_{t-1} + \theta_2 \varepsilon_{t-2} + \dots + \theta_q \varepsilon_{t-q}$$

- Infinite-order moving average process, $MA(\infty)$,

$$y_t = \mu + \sum_{j=0}^{\infty} \theta_j \varepsilon_{t-j} = \mu + \theta_0 \varepsilon_t + \theta_1 \varepsilon_{t-1} + \theta_2 \varepsilon_{t-2} + \dots$$

- With $\theta_0 = 1$

MA models - Higher Order

- After excluding extreme cases,

$$\sum_{j=0}^{\infty} |\theta_j| < \infty$$

- which implies that the coefficients are absolute summable
- Moreover, the process is covariance-stationary when,

$$\sum_{j=0}^{\infty} |\gamma_j| < \infty$$

MA models - Identifying the order

- With a $MA(1)$ process the effect of the shock ε_{t-1} affects the value of y_t
- Hence, the value for the first autocorrelation, $\rho(1)$ should differ from zero but the others would not
- With a $MA(2)$ process the effect of the shocks ε_{t-1} and ε_{t-2} affect the value of y_t
- Hence, the value for the first two autocorrelations, $\rho(1)$ and $\rho(2)$ should differ from zero but the others would not
- This would allow us to use the ACF to identify the order of an $MA(q)$ process

Figure 3: Identifying the order - $MA(1)$, $MA(2)$ & $MA(3)$ process

AR models - Solutions

- Given the $AR(1)$,

$$y_t = \phi y_{t-1} + \varepsilon_t$$

- Relates the value of a variable y at time t , to its previous value at time $(t - 1)$ and a random disturbances ε , also at time t
- Assuming that ε_t is independent and identically distributed white noise, $\varepsilon_t \sim \text{i. i. d. } \mathcal{N}(0, \sigma^2)$
- We showed that if $|\phi| < 1$, the $AR(1)$ is covariance-stationary,

$$\mathbb{E}[y_t] = 0$$

$$\text{var}[y_t] = \frac{\sigma^2}{1 - \phi^2}$$

$$\text{cov}[y_t, y_{t-j}] = \phi^j \text{var}[y_t]$$

- To prove this we use recursive substitution, method of undetermined coefficients, or lag operators

AR models - Recursive Substitution

- Starting at some period of time, j

$$\begin{aligned}y_t &= \phi y_{t-1} + \varepsilon_t \\&= \phi(\phi y_{t-2} + \varepsilon_{t-1}) + \varepsilon_t \\&= \phi^2(\phi y_{t-3} + \varepsilon_{t-2}) + \phi \varepsilon_{t-1} + \varepsilon_t \\&= \vdots \\&= \phi^{j+1} y_{t-(j+1)} + \phi^j \varepsilon_{t-j} + \dots + \phi^2 \varepsilon_{t-2} + \phi \varepsilon_{t-1} + \varepsilon_t\end{aligned}$$

- Explains y as a linear function of the initial value $y_{t-(j+1)}$ and the historical values of ε_t
- If $|\phi| < 1$ and j becomes large, $\phi^{j+1} y_{t-(j+1)} \rightarrow 0$
- Thus, the $AR(1)$ can be expressed as an $MA(\infty)$
- Note that if $|\phi| > 1$ and j becomes large, $\phi^j \rightarrow \infty$
- Hence, the equivalent of an autoregressive random walk is an moving average with coefficients that are not summable

AR models - Lag operators

- Lag operators are particularly useful when dealing with more complex model structures
- The straightforward $AR(1)$ model can be written as,

$$(1 - \phi L) y_t = \varepsilon_t$$

- Such a sequence $\{y_t\}_{t=-\infty}^{\infty}$ is bounded if there exists a finite number k , such that $|y_t| < k$ for all t
- Provided $|\phi| < 1$ and we restrict ourselves to bounded sequences, we can multiply by $(1 - \phi L)^{-1}$ on both sides of the equality (the process is invertible),

$$\begin{aligned}(1 - \phi L)^{-1} (1 - \phi L) y_t &= (1 - \phi L)^{-1} \varepsilon_t \\ y_t &= (1 - \phi L)^{-1} \varepsilon_t\end{aligned}$$

AR models - Lag operators

- Under the assumption that $|\phi| < 1$, we can apply the geometric rule,

$$(1 - \phi L)^{-1} = \lim_{j \rightarrow \infty} \left(1 + \phi L + (\phi L)^2 + \dots + (\phi L)^j \right)$$

- This is based on the expression, $(1 - z)^{-1} = 1 + z + z^2 + z^3 + \dots$, which holds if $|z| < 1$
- Using this we can solve for,

$$y_t = \varepsilon_t + \phi \varepsilon_{t-1} + \phi^2 \varepsilon_{t-2} + \phi^3 \varepsilon_{t-3} + \dots = \sum_{j=0}^{\infty} \phi^j \varepsilon_{t-j}$$

AR models - Lag operators

- This expression could be written as a $MA(\infty)$,

$$y_t = \varepsilon_t + \theta_1 \varepsilon_{t-1} + \theta_2 \varepsilon_{t-2} + \theta_3 \varepsilon_{t-3} + \dots = \sum_{j=0}^{\infty} \theta_j \varepsilon_{t-j}$$

- Therefore, when $|\phi| < 1$,

$$\sum_{j=0}^{\infty} |\theta_j| = \sum_{j=0}^{\infty} |\phi^j|$$

AR models - Unconditional Moments

- The unconditional first-and second-order moments of a stable $AR(1)$ process may be represented by an $MA(\infty)$,
- Where for $y_t = \phi y_{t-1} + \varepsilon_t$,

$$\mathbb{E}[y_t] = \mathbb{E}[\varepsilon_t + \phi\varepsilon_{t-1} + \phi^2\varepsilon_{t-2} + \phi^3\varepsilon_{t-3} + \dots] = 0$$

- The variance is then,

$$\begin{aligned}\gamma[0] &= \text{var}[y_t] = \mathbb{E}[y_t - \mathbb{E}[y_t]]^2 \\ &= \mathbb{E}[\varepsilon_t + \phi\varepsilon_{t-1} + \phi^2\varepsilon_{t-2} + \phi^3\varepsilon_{t-3} + \dots]^2 \\ &= \text{var}[\varepsilon_t] + \phi^2\text{var}[\varepsilon_{t-1}] + \phi^4\text{var}[\varepsilon_{t-2}] + \phi^6\text{var}[\varepsilon_{t-3}] + \dots \\ &= (1 + \phi^2 + \phi^4 + \phi^6 + \dots)\sigma^2 \\ &= \frac{1}{1 - \phi^2}\sigma^2\end{aligned}$$

AR models - Unconditional Moments

- The first order covariance is then,

$$\begin{aligned}\gamma(1) &= \mathbb{E} \left[(y_t - \mathbb{E}[y_t]) (y_{t-1} - \mathbb{E}[y_{t-1}]) \right] \\ &= \mathbb{E} \left[(\varepsilon_t + \phi \varepsilon_{t-1} + \phi^2 \varepsilon_{t-2} + \dots) \times (\varepsilon_{t-1} + \phi \varepsilon_{t-2} + \dots) \right] \\ &= (\phi + \phi^3 + \phi^5 + \dots) \sigma^2 = \phi (1 + \phi^2 + \phi^4 + \dots) \sigma^2 \\ &= \phi \frac{1}{1 - \phi^2} \sigma^2 \\ &= \phi \text{var}[y_t]\end{aligned}$$

- While for $j > 1$ we have,

$$\gamma(j) = \mathbb{E} \left[(y_t - \mathbb{E}[y_t]) (y_{t-j} - \mathbb{E}[y_{t-j}]) \right] = \phi^j \text{var}[y_t]$$

- which proves the result relating to the stationarity of the $AR(1)$ model when $|\phi| < 1$

AR models - Unconditional Moments

- As noted previously the ACF for an $AR(1)$ process coincides with its impulse response function
- where for the ACF of an $AR(1)$ for $j = 1, \dots, J$

$$\rho(0) = \frac{\gamma(0)}{\gamma(0)} = 1, \rho(1) = \frac{\gamma(1)}{\gamma(0)} = \phi, \dots, \rho(j) = \frac{\gamma(j)}{\gamma(0)} = \phi^j$$

- Which equals the dynamic multipliers that may be summarised by the impulse response function

$$\frac{\partial y_t}{\partial \varepsilon_t} = 1, \frac{\partial y_t}{\partial \varepsilon_{t-1}} = \phi, \dots, \frac{\partial y_t}{\partial \varepsilon_{t-j}} = \phi^j$$

Figure 4: Autocorrelation functions for $AR(1)$ processes

AR models - Adding a constant

- To ascertain how the results change after adding a constant,

$$y_t = \mu + \phi y_{t-1} + \varepsilon_t$$

- We can define $v_t = \mu + \varepsilon_t$, such that,

$$\begin{aligned} y_t &= \phi y_{t-1} + v_t \\ y_t &= (1 - \phi L)^{-1} v_t \\ &= \left(\frac{1}{1 - \phi} \right) \mu + \varepsilon_t + \phi \varepsilon_{t-1} + \phi^2 \varepsilon_{t-2} + \dots \end{aligned}$$

- with unconditional first moment,

$$\mathbb{E}[y_t] = \left(\frac{1}{1 - \phi} \right) \mu$$

- which does not depend on time

AR models - Higher order processes

- For higher-order autoregressive processes, things become a bit more complicated, where

$$y_t = \phi_1 y_{t-1} + \phi_2 y_{t-2} + \varepsilon_t$$

- No longer able to consider the value of ϕ_1 alone to determine whether it is stationary
- To complete the process we have to rewrite the $AR(2)$ expression as a first order difference equation

AR models - Higher order processes

- Using a vector, Z_t , which is of dimension (2×1) ,

$$Z_t = \begin{bmatrix} y_t \\ y_{t-1} \end{bmatrix}$$

- With a vector for the errors,

$$v_t = \begin{bmatrix} \varepsilon_t \\ 0 \end{bmatrix}$$

- And the (2×1) matrix for the coefficients,

$$\Gamma = \begin{bmatrix} \phi_1 & \phi_2 \\ 1 & 0 \end{bmatrix}$$

AR models - Higher order processes

- The first-order vector difference equation can be written,

$$Z_t = \Gamma Z_{t-1} + v_t$$

- The matrix Γ is termed the *companion form* matrix of the $AR(2)$ process
- To check for stationarity we can compute the eigenvalues of this matrix
- Moreover, the eigenvalues of Γ are two solutions of x polynomial that satisfy the characteristic equation:

$$x^2 - \phi_1 x - \phi_2 = 0$$

AR models - Higher order processes

- These eigenvalues (m_1 and m_2) must then satisfy $(x - m_1)(x - m_2)$, and can be found from the formula:

$$m_1, m_2 = \frac{\left(\phi_1 \pm \sqrt{\phi_1^2 + 4\phi_2} \right)}{2}$$

- Stationarity requires that the eigenvalues are less than one in absolute value
- In the $AR(2)$ case, one can show that this will be the case if,

$$\begin{aligned}\phi_1 + \phi_2 &< 1 \\ -\phi_1 + \phi_2 &< 1 \\ \phi_2 &> -1\end{aligned}$$

Figure 5: Eigenvalues for difference equation $x^2 - 0.6x - 0.2 = 0$

AR models - Higher order processes

- The $AR(p)$ can then be written as,

$$y_t = \phi_1 y_{t-1} + \phi_2 y_{t-2} + \dots + \phi_p y_{t-p} + \varepsilon_t$$

- Checking for stationarity involves similar calculations
- In this case the Γ matrix will be of the form:

$$\Gamma = \begin{bmatrix} \phi_1 & \phi_2 & \phi_2 & \dots & \phi_{p-1} & \phi_p \\ 1 & 0 & 0 & \dots & 0 & 0 \\ 0 & 1 & 0 & \dots & 0 & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots & \vdots \\ 0 & 0 & 0 & \dots & 1 & 0 \end{bmatrix}$$

- Provided the eigenvalues are less than one in absolute value, (i.e. they lie within the unit circle), the p^{th} order autoregression is stable

AR models - Identify the order of AR(p)

- As in the case of the $MA(q)$ processes one could use the ACF coefficients to identify the order of the $AR(p)$ process
- However, as the $AR(p)$ process passes on the persistence to successive lags so the ACF would not be useful
- As the PACF removes the effects of the persistence that is passed on from intervening lags of the $AR(p)$ process it may be used to identify the order of an $AR(p)$ process

ARMA models

- We can specify an $ARMA(1, 1)$ process as,

$$y_t = \phi y_{t-1} + \varepsilon_t + \theta \varepsilon_{t-1}$$

- Or using the lag polynomials, a general form of an ARMA model is,

$$\phi(L) y_t = \theta(L) \varepsilon_t$$

- Note that the number of lags, (p) and (q) , can differ
- For instance, an $ARMA(2, 1)$ combines an $AR(2)$ with an $MA(1)$:

$$\begin{aligned} (1 - \phi_1 L - \phi_2 L^2) y_t &= (1 + \theta_1 L) \varepsilon_t \\ y_t &= \phi_1 y_{t-1} + \phi_2 y_{t-2} + \varepsilon_t + \theta_1 \varepsilon_{t-1} \end{aligned}$$

ARMA processes

- Whether an $ARMA(p, q)$ process is stationary depends solely on its autoregressive past
- Assume an $ARMA(1, 1)$ and using the lag operator,

$$(1 - \phi L) y_t = (1 + \theta L) \varepsilon_t$$

- Multiplying by $(1 - \phi L)^{-1}$ on both sides,

$$\begin{aligned} y_t &= \frac{(1 + \theta L)}{(1 - \phi L)} \varepsilon_t \\ &= (1 - \phi L)^{-1} \varepsilon_t + (1 - \phi L)^{-1} \theta_1 \varepsilon_{t-1} \end{aligned}$$

ARMA processes

- When $|\phi| < 1$, this can be written as the geometric process,

$$\begin{aligned}y_t &= \sum_{j=0}^{\infty} (\phi L)^j \varepsilon_t + \theta \sum_{j=0}^{\infty} (\phi L)^j \varepsilon_{t-1} \\&= \varepsilon_t + \sum_{j=1}^{\infty} \phi^j \varepsilon_{t-j} + \theta \sum_{j=1}^{\infty} \phi^{j-1} \varepsilon_{t-j} \\&= \varepsilon_t + \sum_{j=1}^{\infty} (\phi^j + \theta \phi^{j-1}) \varepsilon_{t-1}\end{aligned}$$

Figure 6: ACF and PACF for $AR(1)$ with $\phi = 0.5$

Figure 7: ACF and PACF for $MA(1)$ with $\theta = 0.6$

Figure 8: ACF and PACF for $ARMA(1, 1)$ with $\phi = 0.5$ and $\theta = 0.6$

Autocorrelation patterns

- When combining the AR and MA correlation functions the results may be somewhat unclear
- Possibly $ARMA(2, 2)$, $ARMA(1, 2)$, $ARMA(2, 1)$, $ARMA(1, 1)$, $MA(2)$ or $AR(2)$

Seasonal ARMA Models

- In several cases the dependence on the past occurs with a seasonal lag s
- With monthly economic data the behaviour in Jan 2010 may be related to Jan 2011
- Could introduce autoregressive and moving average terms that arise at a seasonal interval
- For example, $ARMA(p, q)_s$ model that takes the form $ARMA(1, 1)_{12}$ would be written as,

$$y_t = \phi y_{t-12} + \varepsilon_t + \theta \varepsilon_{t-12}$$

- Estimation is relatively straightforward

Seasonal ARMA Models - Identification

- The $MA(1)$ with a seasonal ($s = 12$), which could be written as, $y_t = \varepsilon_t + \theta\varepsilon_{t-12}$
- It is easy to verify that

$$\begin{aligned}\gamma(0) &= (1 + \theta^2)\sigma^2 \\ \gamma(12) &= \theta\sigma^2 \\ \gamma(j) &= 0, \text{ for values where } j \neq 12\end{aligned}$$

- The only non-zero autocorrelation, aside from lag zero is, $\rho(12) = \theta/(1 + \theta^2)$

Seasonal ARMA Models - Identification

- Similarly, for the $AR(1)$ model with seasonal ($s = 12$), we could calculate,

$$\begin{aligned}\gamma(0) &= \sigma^2 / (1 - \phi^2) \\ \gamma(12) &= \sigma^2 \phi^k / (1 - \phi^2) \text{ for } k = 1, 2, \dots \\ \gamma(j) &= 0, \text{ for values where } j \neq 12\end{aligned}$$

- Results suggest the PACF from non-seasonal are analogous to the seasonal models

Seasonal ARMA Models - Identification

- Could allow for mixed seasonal models in the general $ARMA(p, q)_s$ framework,

$$y_t = \phi y_{t-12} + \varepsilon_t + \theta \varepsilon_{t-1}$$

- While estimation would be straightforward, the identification of the structural form may be problematic

Box-Jenkins methodology

- In a real world application we would not know the functional form of the underlying data generating process
- The respective parameters in these models would then need to be estimated
- Thereafter, we could assess the model fit
- This procedure is encapsulated in the Box & Jenkins (1979) methodology
 - Identification, Estimation, Diagnostic testing

Box-Jenkins - Identification

- Examine the time plot of the data to
 - detect and correct for outliers, missing values, structural breaks (if possible)
 - detect nonstationary by a pronounced trend or prolonged meander (possibly correct)
- If you are uncertain about the degree of stationarity then perform unit root tests
 - plot ACF and PACF to consider the persistence in the data
 - when ACF quickly returns to zero then there will be no unit root
- Alternatively, if you think that the data represents white noise then use Q -statistic

Box-Jenkins - Identification

- Calculate Q -statistic to test whether a group of autocorrelations is different from zero
- Originally developed by Box-Pierce (1970) better small-sample performance reported by Ljung and Box (1978)

$$Q = T(T + 2) \sum_{k=1}^s \rho_k^2 / (T - k)$$

- If sample value of Q exceeds the critical value χ^2 with s degrees of freedom then at least one value of ρ_k is statistically different from zero at specified significance level

Box-Jenkins - Identification

- Examine the ACF and PACF functions more closely to try to identify the order of a potential $ARMA(p, q)$
- For the ACF and PACF functions that were provided previously we would consider an $ARMA(2, 2)$, an $ARMA(1, 2)$, an $ARMA(2, 1)$ or an $ARMA(1, 1)$
- Would also think about using a $MA(2)$ or an $AR(2)$, but not a $MA(1)$ or an $AR(1)$

Box-Jenkins - Estimation Stage

Fit each of the candidate models and examine the various ϕ_i and θ_i coefficients according to:

- Parsimony:
 - Additional coefficients increase fit but reduce degrees of freedom
 - Parsimonious models often produce better out-of-sample fit
- Stationarity and Invertibility:
 - Distribution theory underlying the use of sample ACF and PACF as approximations for the true DGP assume that y_t is stationary
 - t -statistics and Q -statistics presume that the data is stationary
- Be suspicious if the estimated value of $|\phi_1|$ is close to unity
- Model must be invertible since the ACF and PACF assume that y_t can be approximated by an $AR(1)$ model where $|\phi_1| < 1$

Box-Jenkins - Estimation Stage

- To evaluate the different candidate models consider the goodness-of-fit measures:
 - Look at R^2 and average of the residual sum of squares
 - AIC and BIC are more suitable criteria since they weigh-up parsimony and "goodness-of-fit"
 - Smaller values of a AIC are better (or where $AIC < 0$, choose the model with the most negative statistic)

Box-Jenkins estimation - AIC & BIC

- Adding additional lags will reduce the sum of squares of the estimated residuals (and will lead to a higher R^2)
- But you will also loose degrees of freedom (which may be essential)
- Akaike and Bayesian Information Criteria test for goodness of fit, while prizing parsimony

$$AIC = \log \hat{\sigma}_k^2 + \frac{T + 2k}{T}$$

$$BIC = \log \hat{\sigma}_k^2 + \frac{k \log T}{T}$$

- where k = number of estimated parameters and T = nobs
- $\hat{\sigma}_k^2 = \frac{SSR_k}{T}$ is the variance of the residual sum of squares

Box-Jenkins - Estimation Stage

- Make sure T is fixed, when comparing $AR(1)$ & $AR(2)$
- Including a parameter must decrease SSR_k if AIC or BIC is to decrease
- Since $\log T$ is greater than 2, BIC likes more parsimonious models

Box-Jenkins - Diagnostic checking

- Plot the residuals to look for outliers or periods where the model does not fit the data
- Construct ACF and PACF of the residuals
- Serial correlation in the residuals implies that a systematic movement in the y_t sequence is not accounted for by the $ARMA(p, q)$ coefficients
 - Those models should be eliminated and re-estimated
 - Use Q -statistic to determine whether any or all of the ACF or PACF coefficients are significant
- When applying the Q -statistic to the residuals of an $ARMA(p, q)$ model use χ^2 with $s - p - q$ degrees of freedom
 - Ensure that the standard errors for the coefficient estimates are appropriate, if not re-estimate model

Box-Jenkins - Diagnostics & Forecasts

- If possible fit $ARMA(p, q)$ models to subsamples - stability of DGP

$$F = \frac{(RSS - RSS_1 - RSS_2)/n}{(RSS_1 + RSS_2)/(T - 2k)}$$

- where k is the number of parameters, i.e. $p + q + 1$ (with constant)
- If all the coefficients are equal ($RSS_1 + RSS_2$) should equal RSS and $F = 0$
- You could then use the model for forecasting y_{T+1}, y_{T+1}, \dots for out-of-sample comparison

Structural Breaks - Chow's Breakpoint

- Model two sub-samples of the data and see whether there are significant differences in the parameters
- Test whether the null hypothesis of "no structural change" holds after constructing a F test statistic for the parameters
- Could construct a model for changes at date τ

$$y_t = x_t^\top \beta_t + \varepsilon_t$$

- where

$$\beta_t = \begin{cases} \beta & t \leq \tau \\ \beta + \delta & t > \tau \end{cases}$$

- Or alternatively we could test for change in all the model parameters with an F test

Structural Breaks - Chow's Breakpoint

- Major drawback is that the change point must be known *a priori*
- Must ensure that each sub-sample has at least as many observations as the number of estimated parameters

Structural Breaks - Quandt LR Test

- Extension of the Chow test where an F test statistic is calculated for all potential breakpoints within an interval $[\underline{i}, \bar{i}]$
- Reject the null hypothesis of no structural change if the absolute value of any of the test statistics are too large
- Takes the form of a sup F test
- Asymptotic properties of this statistic are non-standard so use those that are referenced in the notes

Figure 9: Quandt Likelihood Ratio Test - Breakpoint at observation 100 with $n = 200$ and $p = 0.00$

Structural Breaks - CUSUM Test

- CUSUM test is based on the cumulative sum of the recursive residuals
- Plot the cumulative sum together with the 5% critical boundaries
- If the cumulative sum breaks either of the two boundaries there is parameter instability and a possible structural break
- Need to specify the model *a priori* to obtain the residuals

Figure 10: CUSUM Test - Breakpoints for change in coeffic

Conclusion

- Relatively simple ARMA models can be used to describe stationary univariate time series
- Easy to estimate and use of the straightforward Box & Jenkins method can identify possible functional forms for the underlying data generating process
- It is possible to test the data generating process for structural breaks